

CS1952Q HW

David Heffren

February 2023

1 Introduction

In this assignment, to find X and Y minimizing $f(X, Y) = \sum_{(i,j) \in \Omega} (M_{ij} - XY_{ij}^T)^2$, I used batch gradient descent in python (considering a single row a batch of 1), alternating between X and Y .

I did this only using the numpy and cupy packages in python. I split the given examples into 3 groups: training, validation, and test, with a split of .8/.1/.1.

To deal with the issue that we didn't have all values of M , and thus wanted to ignore some terms in the sum (but to avoid using slow for loops in python), I created a matrix M with a value of 0 if we didn't have the pair (i,j) in the training data, and a boolean matrix B with 1 if we have the pair and 0 if we don't.

With this, I could rewrite the sums in matrix form, and speed up the calculations with numpy, avoiding for loops. Thus, the loss function became:

$$f(X, Y) = \| (M - XY^T) * B \|^2 = \sum_{i,j} (M - XY^T * B)_{ij}^2$$

where $*$ is the elementwise product.

For the gradient of X, Y I calculated it in matrix form as:

$$\begin{aligned}\frac{\partial f}{\partial X} &= -2(M - XY^T * B)Y \\ \frac{\partial f}{\partial Y} &= -2(M - XY^T * B)^T X\end{aligned}$$

Note that originally I just took the gradient of the WHOLE matrix X and Y at once, and updated them that way, but this caused the problem that the gradients would explode and I'd end up with nans. So, I first implemented it by taking the gradient of each row individually (calculating the entire gradient then taking a row of that was way too inefficient), and updating X and Y row by row. However, this was too slow, so I instead updated the gradient via batches of X and Y , updating bs adjacent rows of X , then bs adjacent rows of Y . With this, I got about 1 second per epoch.

I tried a couple of different initializations for X and Y , as I first just made each element $N(0,1)$ for both. This worked well enough. Then, I tried the rank r approximation of the SVD of M (with zeros).

To choose r , I essentially used trial and error. I wanted to pick the minimum r value s.t it achieved very good results. When r was too large, it wouldn't learn much. I ended up settling on an r of 6, although $r=4$ worked well, and took many fewer epochs. Anything higher than 10 didn't work well for me at all.

To calculate the loss itself, I used the above stated loss function/number of examples, in order to match what the graders use.

With $r = 6$, $lr = .00004$, $epochs = 500$, $bs = 10$, I managed to get pretty good results:

train: .642,

val: .6645,

test: .664.

As you can see, this .664 is well below the threshold .8 for the assignment. Since this test data was sampled randomly from the list of all given data, it should be an accurate predictor of the randomly sampled data used by the graders.

For $r = 4$, my best results were train: .7265, val: .7419, test: .7432, with learning rate .00005, $bs = 10$, $epochs = 100$. Note that this learning rate became too large towards the end, as the training and validation loss began to increase.

I wanted to try a learning rate with exponential decay.

Note I originally found decent values for r on the smaller dataset, because it was loads faster to run (Especially before I optimized). It turns out the good values for r doesn't depend on the size of the dataset. However, the learning rate definitely does.

I added a regularizer which penalizes XY^T which go above 5 or below .5, in order to incentivize the results are in the right range. This decreases the test loss to around .6.